

# **LIKWID:**



## **Lightweight performance tools**

**J. Treibig**

**RRZE, University Erlangen**

**26.9.2011**

For high efficiency hardware aware programming is required.



Multicore architectures add complex topologies on the thread and memory level.

- Lightweight command line tools for Linux
- Help to face the challenges without getting in the way
- Focus on X86 architecture
- Philosophy:
  - Simple
  - Efficient
  - Portable
  - Extensible



Open source project (GPL v2):

<http://code.google.com/p/likwid/>




- **Question: There is tool XY? They can do the same thing.**
- **Possible answers:**
  - LIKWID has an unique feature set
  - LIKWID has NO external dependencies
  - LIKWID is easy to build and setup
  - LIKWID is just COOL (OK this is biased)

**If you are still not convinced:**

**It is always good to have some competition.**

**Even in Open Source tools.**

**So try it and make your own opinion what suits your needs best.**



## **Scenario 1: Hardware performance monitoring and Node performance characterisation**

likwid-perfctr  
likwid-perfscope  
likwid-bench



- **A coarse overview of hardware performance monitoring data is often sufficient**

- **likwid-perfctr** (similar to “perfex” on IRIX, “hpmcount” on AIX, “lipfpm” on Linux/Altix, “craypat” on Cray systems)
- Simple end-to-end measurement of hardware performance metrics
- Operating modes:
  - Wrapper
  - Stethoscope
  - Timeline
  - Marker API

- Preconfigured and extensible metric groups, list with **likwid-perfctr -a**



BRANCH: Branch prediction miss rate/ratio  
CACHE: Data cache miss rate/ratio  
CLOCK: Clock of cores  
DATA: Load to store ratio  
FLOPS\_DP: Double Precision MFlops/s  
FLOPS\_SP: Single Precision MFlops/s  
FLOPS\_X87: X87 MFlops/s  
L2: L2 cache bandwidth in MBytes/s  
L2CACHE: L2 cache miss rate/ratio  
L3: L3 cache bandwidth in MBytes/s  
L3CACHE: L3 cache miss rate/ratio  
MEM: Main memory bandwidth in MBytes/s  
TLB: TLB miss rate/ratio



```
$ env OMP_NUM_THREADS=4 likwid-perfctr -C N:0-3 -t intel -g FLOPS_DP ./stream.exe
```

```
-----
CPU type:      Intel Core Lynnfield processor
CPU clock:     2.93 GHz
-----
```

```
Measuring group FLOPS_DP
```

```
YOUR PROGRAM OUTPUT
```

Event	core 0	core 1	core 2	core 3
INSTR_RETIRED_ANY	1.97463e+08	2.31001e+08	2.30963e+08	2.31885e+08
CPU_CLK_UNHALTED_CORE	9.56999e+08	9.58401e+08	9.58637e+08	9.57338e+08
FP_COMP_OPS_EXE_SSE_FP_PACKED	4.00294e+07	3.08927e+07	3.08866e+07	3.08904e+07
FP_COMP_OPS_EXE_SSE_FP_SCALAR	882	0	0	0
FP_COMP_OPS_EXE_SSE_SINGLE_PRECISION	0	0	0	0
FP_COMP_OPS_EXE_SSE_DOUBLE_PRECISION	4.00303e+07	3.08927e+07	3.08866e+07	3.08904e+07

Always  
measured

Configured metrics  
(this group)

Metric	core 0	core 1	core 2	core 3
Runtime [s]	0.326242	0.32672	0.326801	0.326358
CPI	4.84647	4.14891	4.15061	4.12849
DP MFlops/s (DP assumed)	245.399	189.108	189.024	189.304
Packed MUOPS/s	122.698	94.554	94.5121	94.6519
Scalar MUOPS/s	0.00270351	0	0	0
SP MUOPS/s	0	0	0	0
DP MUOPS/s	122.701	94.554	94.5121	94.6519

Derived  
metrics



- **likwid-perfctr measures on core base and has no notion what runs on the cores**

**This enables to listen on what currently happens without any overhead:**



```
likwid-perfctr -c N:0-11 -g FLOPS_DP sleep 10
```

- **It can be used as cluster/server monitoring tool**
- **A frequent use is to measure a certain part of a long running parallel application from outside**

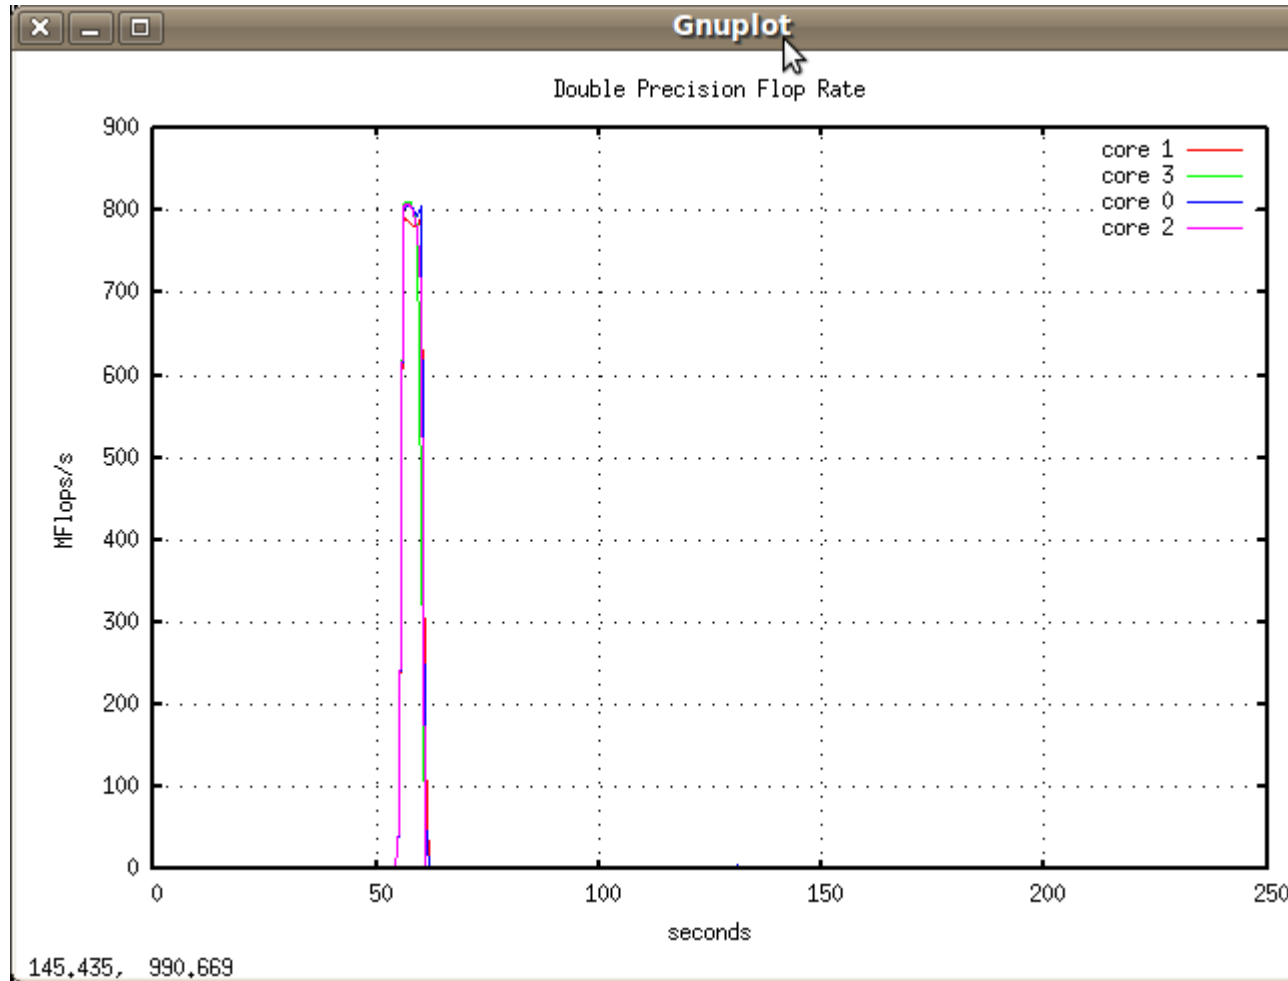


- likwid-perfctr supports time resolved measurements of full node:

```
likwid-perfctr -c N:0-11 -g MEM -d 50ms > out.txt
```



```
likwid-perfscope -group FLOPS_DP -cores 0-3
```





- To measure only parts of an application a marker API is available.
- The API only turns counters on/off. The configuration of the counters is still done by likwid-perfctr application.
- Multiple named regions can be measured
- Results on multiple calls are accumulated
- Inclusive and overlapping Regions are allowed

```
likwid_markerInit(); // must be called from serial region
```

```
likwid_markerStartRegion("Compute");
```

```
...
```

```
likwid_markerStopRegion("Compute");
```

```
likwid_markerStartRegion("postprocess");
```

```
...
```

```
likwid_markerStopRegion("postprocess");
```

```
likwid_markerClose(); // must be called from serial region
```

SHORT PSTI

### EVENTSET

```
FIXC0 INSTR_RETIRED_ANY
FIXC1 CPU_CLK_UNHALTED_CORE
FIXC2 CPU_CLK_UNHALTED_REF
PMC0  FP_COMP_OPS_EXE_SSE_FP_PACKED
PMC1  FP_COMP_OPS_EXE_SSE_FP_SCALAR
PMC2  FP_COMP_OPS_EXE_SSE_SINGLE_PRECISION
PMC3  FP_COMP_OPS_EXE_SSE_DOUBLE_PRECISION
UPMC0 UNC_QMC_NORMAL_READS_ANY
UPMC1 UNC_QMC_WRITES_FULL_ANY
UPMC2 UNC_QHL_REQUESTS_REMOTE_READS
UPMC3 UNC_QHL_REQUESTS_LOCAL_READS
```

### METRICS

```
Runtime [s] FIXC1*inverseClock
CPI      FIXC1/FIXC0
Clock [MHz] 1.E-06*(FIXC1/FIXC2)/inverseClock
DP MFlops/s (DP assumed) 1.0E-06*(PMC0*2.0+PMC1)/time
Packed MUOPS/s 1.0E-06*PMC0/time
Scalar MUOPS/s 1.0E-06*PMC1/time
SP MUOPS/s 1.0E-06*PMC2/time
DP MUOPS/s 1.0E-06*PMC3/time
Memory bandwidth [MBytes/s] 1.0E-06*(UPMC0+UPMC1)*64/time;
Remote Read BW [MBytes/s] 1.0E-06*(UPMC2)*64/time;
```

### LONG

Formula:

```
DP MFlops/s = (FP_COMP_OPS_EXE_SSE_FP_PACKED*2 + FP_COMP_OPS_EXE_SSE_FP_SCALAR)/ runtime.
```

- Groups are architecture specific
- They are defined in simple text files
- During recompile the code is generated
- likwid-perfctr -a outputs list of groups
- For every group an extensive documentation is available



### Likwid supports to specify an output file with placeholder for:

- %j - PBS\_JOBID taken from environment
- %r - MPI Rank as specified by newer Intel MPI versions
- %h - hostname
- %p - process pid

### Example:

```
likwid-perfctr -c L:0 -g FLOPS_DP -o test_%h_%p.txt ./a.out
```

Depending on the file suffix a converter script is called:

- txt Direct output without conversion
- csv Convert to comma separated values format
- xml Convert to xml format

Useful for integration in other tool chains or automated frameworks.



- **Implemented completely in user space (uses msr kernel module)**
- **For security sensitive environments a small proxy application managing a controlled access to the msr device files is available**
- **Supported processors:**
  - Intel Core 2
  - Intel Nehalem /Westmere (all variants) supporting Uncore events
  - Intel NehalemEX/WestmereEX (without Uncore)
  - Intel Sandy Bridge
  - Intel Atom
  - AMD K8/K10
- **likwid-perfctr allows to specify arbitrary event sets on the command line:**

```
likwid-perfctr -c 0-12 -g
```

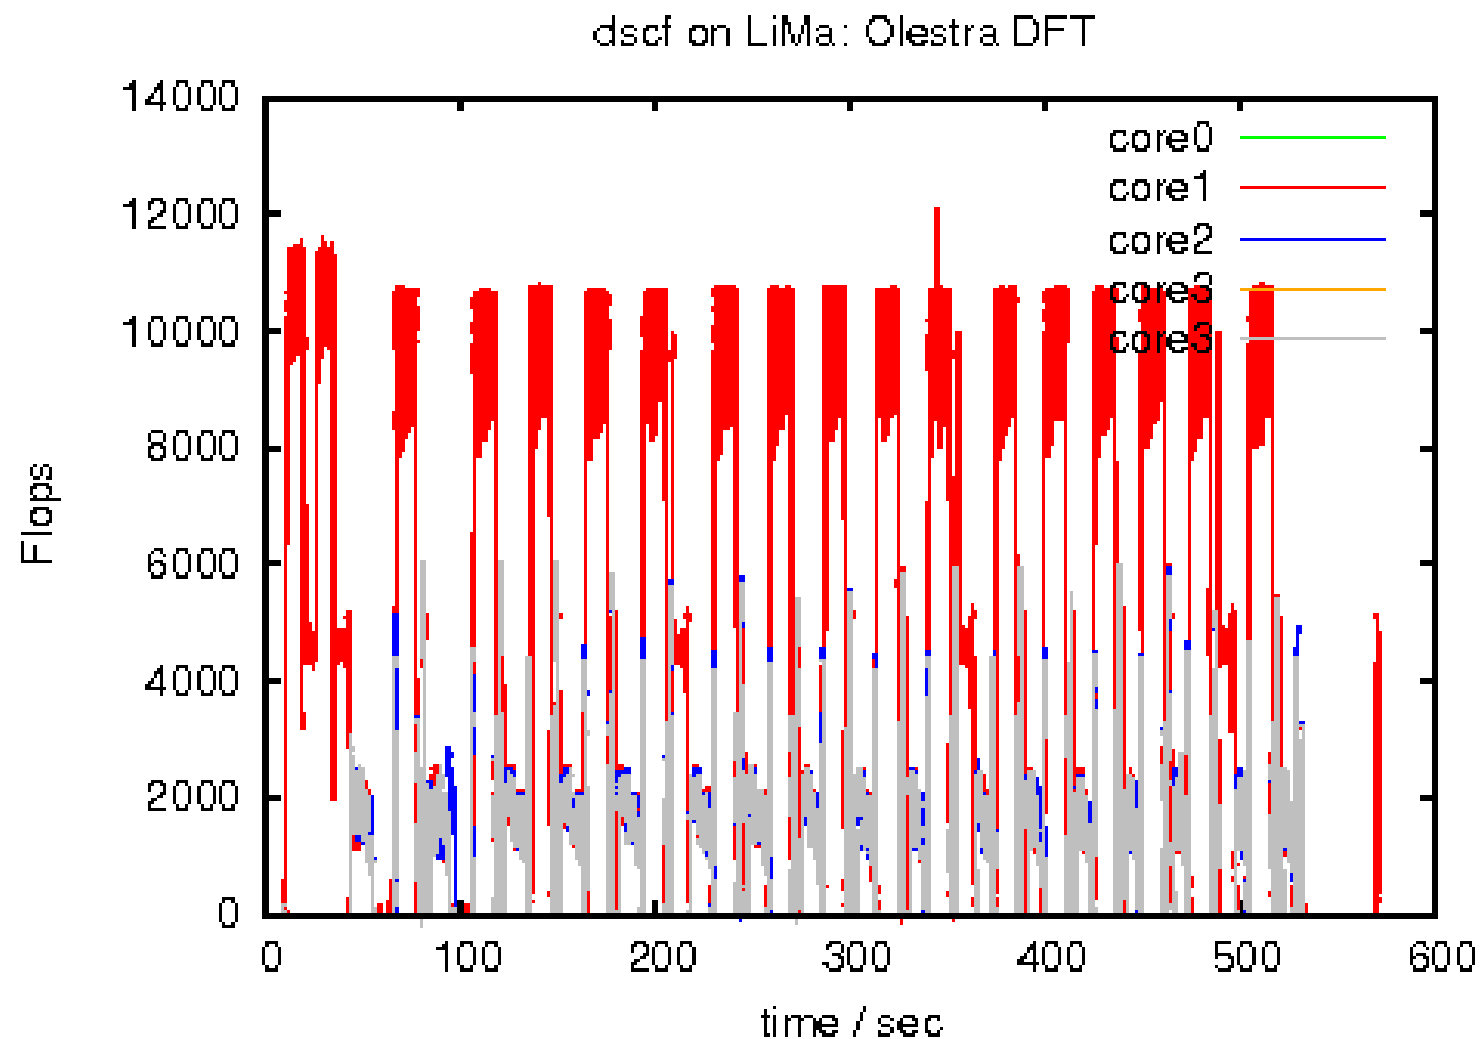
```
INSTR_RETIRED_ANY:FIXC0,CPU_CLK_UNHALTED_CORE:FIXC1,FP_COMP  
_OPS_EXE_SSE_FP_PACKED:PMC0,UNC_L3_LINES_IN_ANY:UPMC0  
sleep 10
```



- **likwid-perfctr can be used with MPI if processes are pinned**
- **For hybrid usage a taskset cpuset must be established**
- **To distinguish the output it can be written to separate files**

```
likwid-perfctr -c L:0 -g FLOPS_DP -o myTag_%r_%h.txt ./app
```

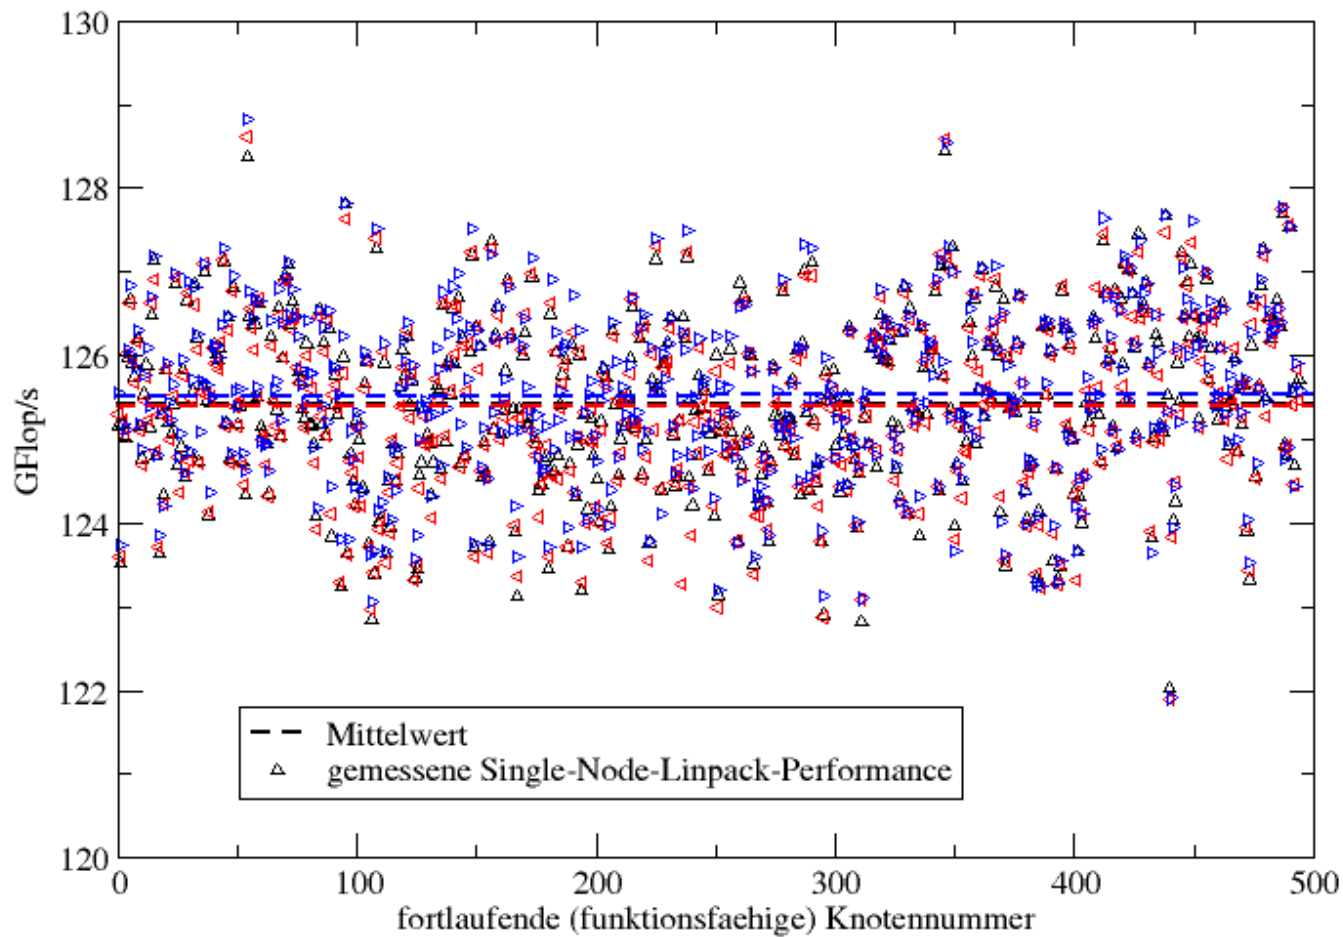
- **There are efforts to add likwid support in Scalaska (and Vampir ?)**
- **likwid-mpirun will have support for perfctr in the future**
- **Well suited to be integrated into batch job system**





## Single-Node Linpack (N=50000, OMP=12)


2010-10-08 / 2x 2010-10-10



## Turbo mode causes volatile performance values


single-node linpack on L0943: 128.4831 GFlop/s

2010-10-12; likwid-perfctr -c 0-11 -g CLOCK -d 5



single-node linpack on L1342: 121.7198 GFlop/s

2010-10-12; likwid-perfctr -c 0-11 -g CLOCK -d 5





- **Full Uncore support for Intel EX type processors**

```
WBOX4  UNCORE_CYCLES
MBOXA0  FVC_EV0_BBOX_CMDS_READS
MBOXB0  FVC_EV0_BBOX_CMDS_READS
BBOXA1  IMT_INSERTS_WR
BBOXB1  IMT_INSERTS_WR
RBOXA0  NEW_PACKETS_RECV_PORT0_IPERF0_ANY_DRS
RBOXA1  NEW_PACKETS_RECV_PORT1_IPERF0_ANY_DRS
RBOXB0  NEW_PACKETS_RECV_PORT4_IPERF0_ANY_DRS
RBOXB1  NEW_PACKETS_RECV_PORT5_IPERF0_ANY_DRS
```

- **Full support for AMD Interlagos**
- **Support for IBM Power 7**
- **Multiplexing**
- **Extension of likwid-perfscope**
- **Porting to Windows and MacOSX**



- **To know the performance properties of a machine is essential for any optimization effort**
- **Microbenchmarking is an important tool to gain this information**
- **Extensible, flexible benchmarking framework**
- **Rapid development of low level kernels**
- **Already includes many ready to use threaded benchmark kernels**
- **Benchmarking runtime cares for:**
  - Thread management and placement
  - Data allocation and NUMA aware initialization
  - Timing and result presentation



- Implement micro benchmark in abstract assembly
- Add meta information
- The benchmark file is automatically converted, compiled and added to the benchmark application

```
$likwid-bench -t clcopy -g 1 -i 1000 -w S0:1MB:2
```

```
$likwid-bench -t load -g 2 -i 100 -w S1:1GB -w S0:1GB-0:S1,1:S0
```

STREAMS 2

TYPE DOUBLE

FLOPS 0

BYTES 16


LOOP 32

```
movaps    FPR1, [STR0 + GPR1 * 8 ]
movaps    FPR2, [STR0 + GPR1 * 8 + 64 ]
movaps    FPR3, [STR0 + GPR1 * 8 + 128 ]
movaps    FPR4, [STR0 + GPR1 * 8 + 192 ]
movaps    [STR1 + GPR1 * 8 ], FPR1
movaps    [STR1 + GPR1 * 8 + 64 ], FPR2
movaps    [STR1 + GPR1 * 8 + 128 ], FPR3
movaps    [STR1 + GPR1 * 8 + 192 ], FPR4
```

- Testcase memcopy (without NT stores)
- tested with likwid-bench on dual socket Nehalem node
- Total bandwidth as reported by likwid-bench
- Bandwidths measured from outside with

```
likwid-perfctr -c S0:0@S1:0 -g MEM
```

**11601 MB/s**




```
likwid-bench -g 1 -i 1000 -t copy -w S0:500MB:4
```

# Detecting NUMA problems 2



- Thread group (4 threads) on socket 0
- Store stream to socket 0 memory
- Load stream from **socket 1** memory

```
likwid-bench -g 1 -i 1000 -t copy -w S0:500MB:4-0:S1,1:S0
```




# Detecting NUMA problems 3



- 2 thread group (2x4 threads) on socket 0/1
- Correct NUMA placement (first touch)

```
likwid-bench -g 2 -i 1000 -t copy -w S0:500MB:4  
-w S1:500MB:4
```


**23147 MB/s**







- 2 thread group (2x4 threads) on socket 0/1
- Common problem: memory placed on one socket only

```
likwid-bench -g 2 -i 1000 -t copy -w S0:500MB:4  
-w S1:500MB:4-0:S0,1:S0
```








Bandwidth map created with likwid-bench. All cores used in one NUMA domain, memory is placed in a different NUMA domain.

Test case: simple copy  $A(:,) = B(:,)$ , large arrays



# AMD Magny Cours 2-socket system


*4 chips, two sockets*



# AMD Magny Cours 4-socket system

*Topology at its best?*





## **Scenario 2: Dealing with node topology and thread affinity**

likwid-topology

likwid-pin

likwid-mpirun



- **Node information is usually scattered in various places**
- **likwid-topology provides all information in a single reliable source**
- **All information is based on cpuid directly**
- **Features:**
  - Thread topology
  - Cache topology
  - NUMA topology
  - Detailed cache parameters (-c command line switch)
  - Processor clock (measured)
  - ASCII art output (-g command line switch)

# Usage: likwid-topology



```
-----
CPU type:      Intel Core Westmere processor
*****
Sockets:      2
Cores per socket: 6
Threads per core: 2
-----
```

```
-----
HWThread      Thread      Core      Socket
0              0           0           0
1              0           1           0
2              0           2           0
-----
```

```
-----
Socket 0: ( 0 12 1 13 2 14 3 15 4 16 5 17 )
Socket 1: ( 6 18 7 19 8 20 9 21 10 22 11 23 )
-----
```

## Cache Topology

```
-----
Level:      3
Size:      12 MB
Type:      Unified cache
Associativity: 16
Number of sets: 12288
Cache line size: 64
Non Inclusive cache
Shared among 12 threads
Cache groups: ( 0 12 1 13 2 14 3 15 4 16 5 17 ) ( 6 18 7 19 8
-----
```

## NUMA Topology


```
-----
NUMA domains: 2
-----
```

### Domain 0:

```
Processors: 0 1 2 3 4 5 12 13 14 15 16 17
Memory: 11615.9 MB free of total 12276.3 MB
-----
```

### Domain 1:




```
Processors: 6 7 8 9 10 11 18 19 20 21 22 23
Memory: 12013.9 MB free of total 12288 MB
-----
```



Information can also be queried in library.  
NUMA information extracted from Linux sys fs.

# Example: STREAM benchmark on 12-core Intel Westmere:

## *Anarchy vs. thread pinning*





There are several reasons for caring about affinity:

- Eliminating performance variation
- Making use of architectural features
- Avoiding resource contention





- Core numbering may vary from system to system even with identical hardware
  - likwid-topology delivers this information, which can then be fed into likwid-pin
- Alternatively, likwid-pin can abstract this variation and provide a purely **logical** numbering (**physical cores first**)



- Across all cores in the node:  
`OMP_NUM_THREADS=8 likwid-pin -c N:0-7 ./a.out`
- Across the cores in each socket and across sockets in each node:  
`OMP_NUM_THREADS=8 likwid-pin -c S0:0-3@S1:0-3 ./a.out`

### Possible unit prefixes

**N**

**node**

**S**


**socket**

**M**


**NUMA domain**

**C**

**outer level cache group**



- ... and: Logical numbering inside a pre-existing cpuset:



- `OMP_NUM_THREADS=4 likwid-pin -c L:0-3 ./a.out`



- Pins process and threads to specific cores without touching code
- Directly supports pthreads, gcc OpenMP, Intel OpenMP
- Allows user to specify skip mask (hybrid OpenMP/MPI)
- Can also be used as replacement for taskset

### Supported usage modes:

- Physical numbering: `likwid-pin -c 0,2,5-8`
- Logical numbering (node): `likwid-pin -c N:3-7`
- Logical numbering (socket): `likwid-pin -c S0:0,2@S2:0-3`
- Logical numbering (NUMA): `likwid-pin -c M0:1-3@M2:1-3`
- Logical numbering (cpuset): `likwid-pin -c L:3-7`


**All logical numberings have physical cores first.**




- Effective improvement without any code change possible

- Memory policy is set to interleave with likwid-pin:

```
likwid-pin -c N:0-7 -i likwid-bench -g 2 -i 1000 -t  
copy -w S0:500MB:4 -w S1:500MB:4-0:S0,1:S0
```



- Uses LD\_PRELOAD for pthread\_create
- Wrapper application controls library through environment variables
- Upon creation threads corresponding to bit position in a skip mask are not pinned





- On the long run a unified standard is needed
- Till then likwid provides a portable solution
- The examples here are for Intel MPI/OpenMP programs, but are also applicable to other threading models

**At the moment the following issues arise**

- Does the OpenMP implementation use a shepherd thread?
- Does the MPI implementation use a shepherd thread?
- Which threadIDs need to be skipped?

**Pure MPI:**


```
likwid-mpirun -np 16 -NperDomain S:2 ./a.out
```

**Hybrid:**

```
likwid-mpirun -np 16 -pin S0:0,1_S1:0,1 ./a.out
```



```
OMP_NUM_THREADS=12 likwid-mpirun -np 2 -pin N:0-11 ./a.out
```



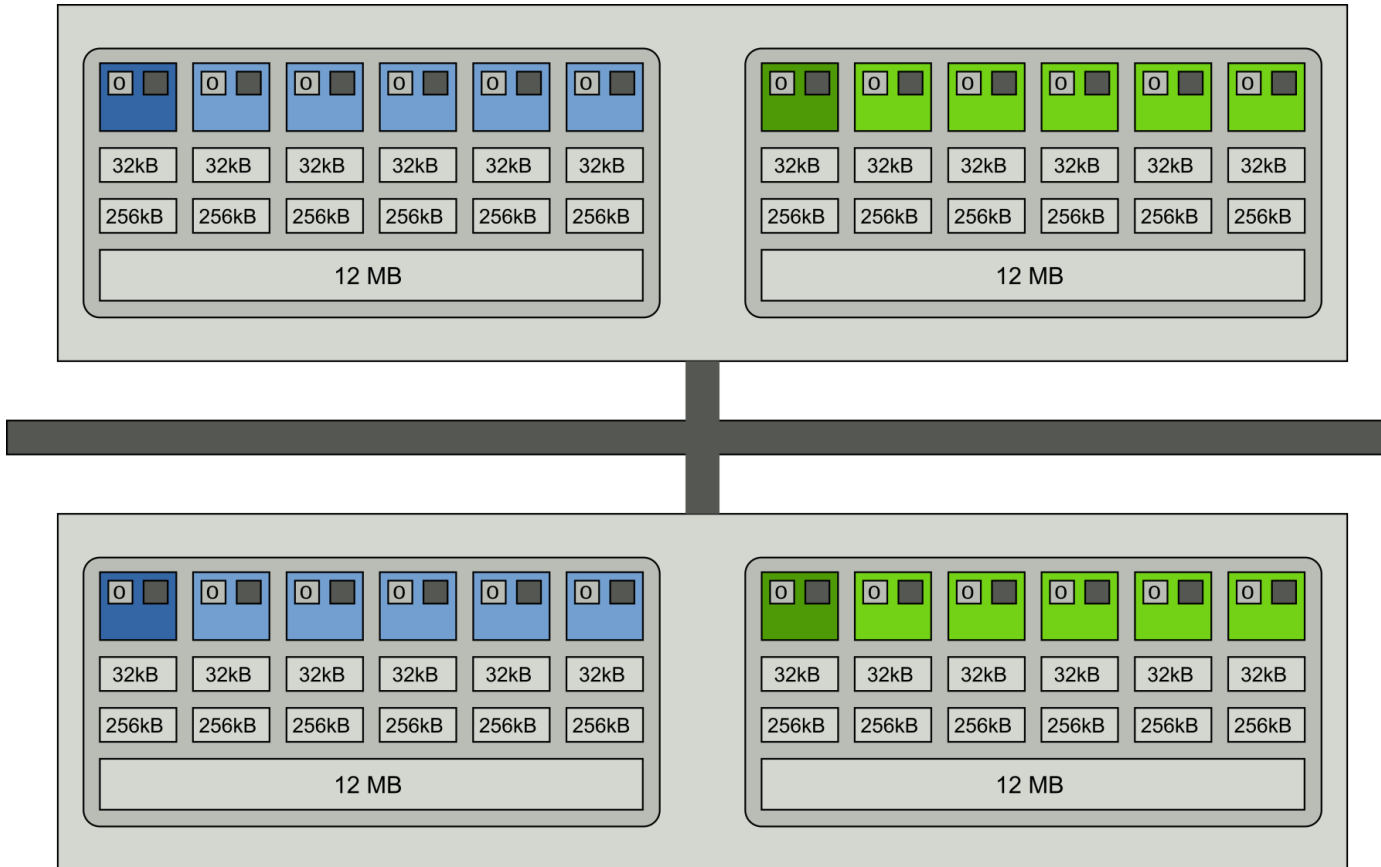
## Intel MPI+compiler:

```
OMP_NUM_THREADS=12 mpirun -ppn 1 -np 2 -env KMP_AFFINITY scatter ./a.out
```





```
OMP_NUM_THREADS=6  likwid-mpirun -np 4 -pin S0:0-5_S1:0-5 ./a.out
```

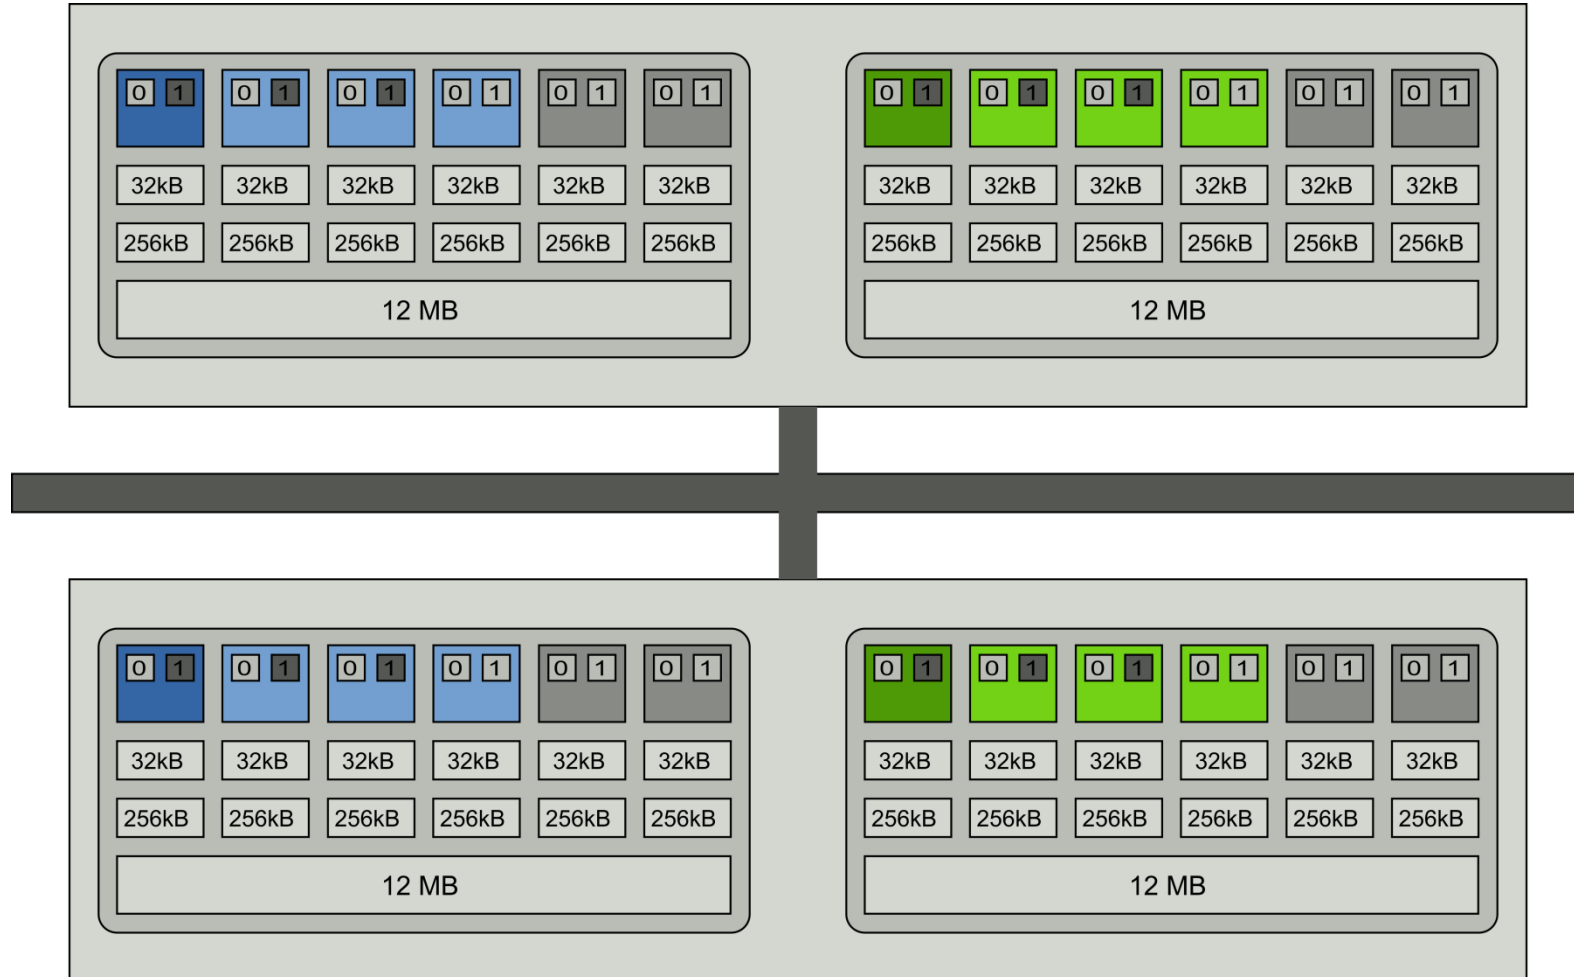


### Intel MPI+compiler:

```
OMP_NUM_THREADS=6 mpirun -ppn 2 -np 4 \  
-env I_MPI_PIN_DOMAIN socket -env KMP_AFFINITY scatter ./a.out
```



```
OMP_NUM_THREADS=5 likwid-mpirun -np 4 -pin S0:0-3,9_S1:0-3,9 ./a.out
```





- **Implements Intel RAPL interface (Sandy Bridge)**
- **RAPL (Running average power limit)**

CPU name: Intel Core SandyBridge processor

CPU clock: 3.49 GHz

Thermal Spec Power: 95 Watts

Minimum Power: 20 Watts

Maximum Power: 95 Watts

Maximum Time Window: 0.15625 micro sec

Energy consumed: 126.597 Joules

Power consumed: 63.2983 Watts

Test case	Power
4 cores, plain C	45.25 Watt
4 cores, SSE	58.74 Watt
4 cores (SMT), SSE	65.71 Watt
4 cores (SMT), AVX	77.85 Watt



- **Automated test suite for**
  - Feature tests on all supported architectures
  - Quantified overhead
  - Automatic validation and deviation tables for performance groups
- **Improved Documentation (WIKI pages)**
- **Adopt likwid for library use**
- **Document usage of likwid as library**
- **Test likwid-mpirun with all relevant MPI implementations**



**Thank you for your attention!**

**Any Questions?**